

Lecture 21 [30.03.2022]

Optimizing Cache Memory Access Time

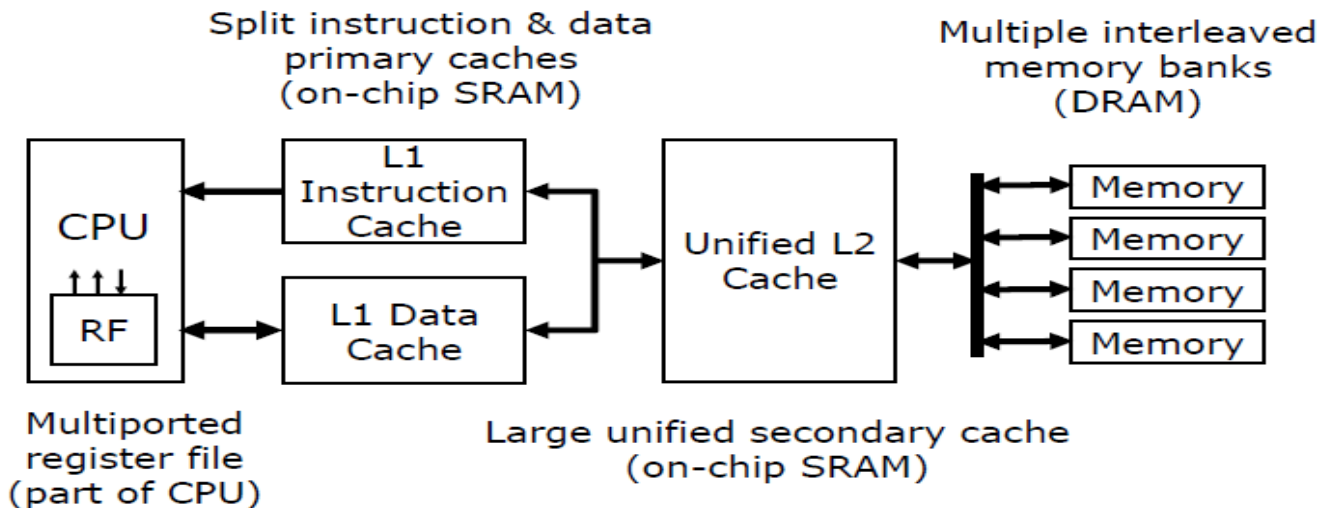
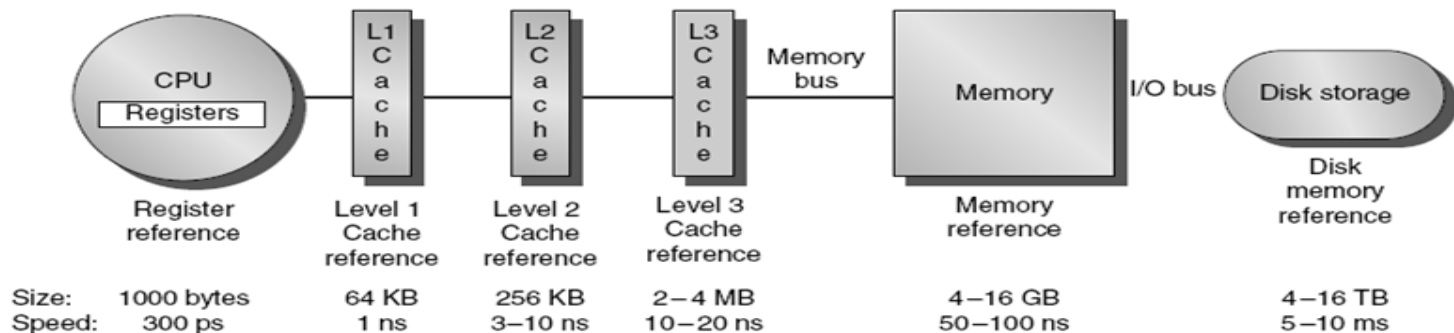


Dr. John Jose

Associate Professor

**Department of Computer Science & Engineering
Indian Institute of Technology Guwahati, Assam.**

Memory Hierarchy




Four cache memory design choices

- ❖ Where can a block be placed in the cache?
 - **Block Placement**
- ❖ How is a block found if it is in the upper level?
 - **Block Identification**
- ❖ Which block should be replaced on a miss?
 - **Block Replacement**
- ❖ What happens on a write?
 - **Write Strategy**

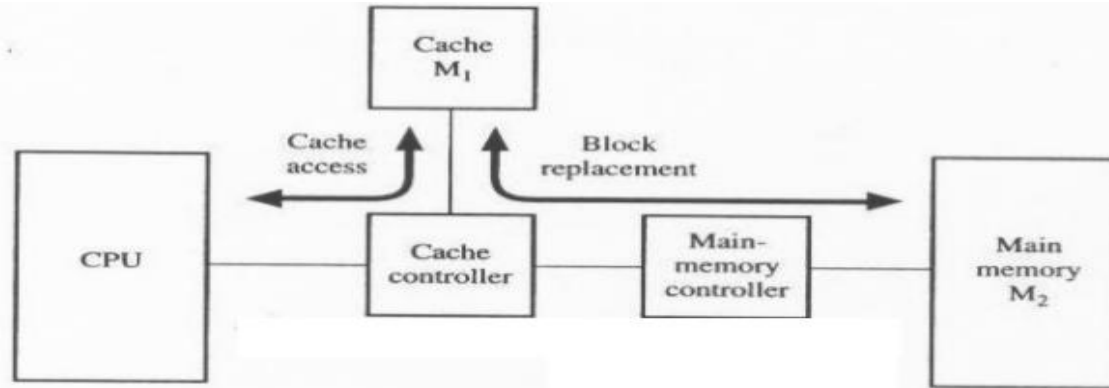
Look-aside vs Look through caches

- ❖ **Look-aside cache:** Request from processor goes to cache and main memory in parallel
- ❖ Cache and main memory both see the bus cycle
- ❖ On cache hit → processor loaded from cache, bus cycle terminates; On cache miss: processor & cache loaded from memory in parallel



Look-aside vs Look through caches

- ❖ **Look-through cache:** Cache checked first when processor requests data from memory
- ❖ On hit → data loaded from cache: On miss → cache loaded from memory, then processor loaded from cache





Write strategy

- ❖ **Write Hits → Write through vs Write back**
- ❖ **Write Miss → Write allocate vs No-Write allocate**
- ❖ **Write through:** The information is written to both the block in the cache and to the main memory
- ❖ Read misses do not need to write back evicted line contents
- ❖ **Write back:** The information is written only to the block in the cache. The modified cache block is written to main memory only when it is replaced.
- ❖ Have to maintain whether block clean or dirty. No extra work on repeated writes; only the latest value on eviction gets updated in main memory.

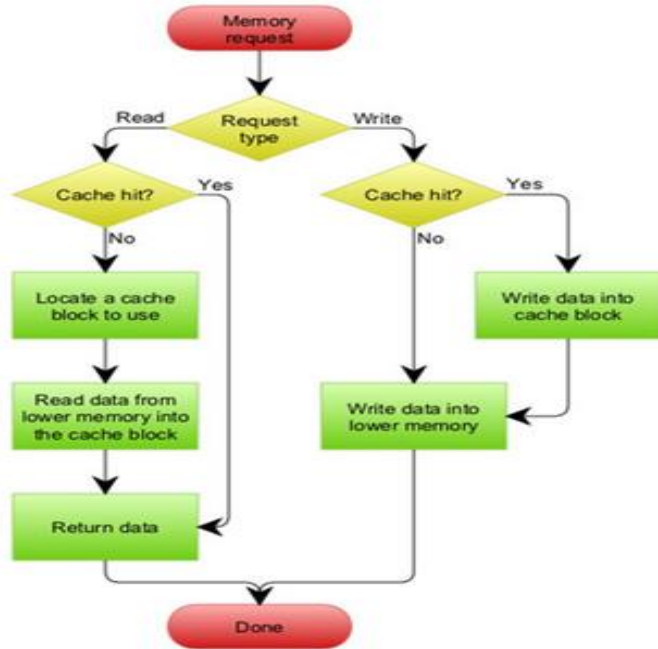
Write strategy

- ❖ **Write allocate:** The block is loaded into cache on a write miss.
- ❖ Used along with write back caches
- ❖ **No-Write allocate:** The block is modified in the main memory but not in cache
- ❖ Used along with write through caches

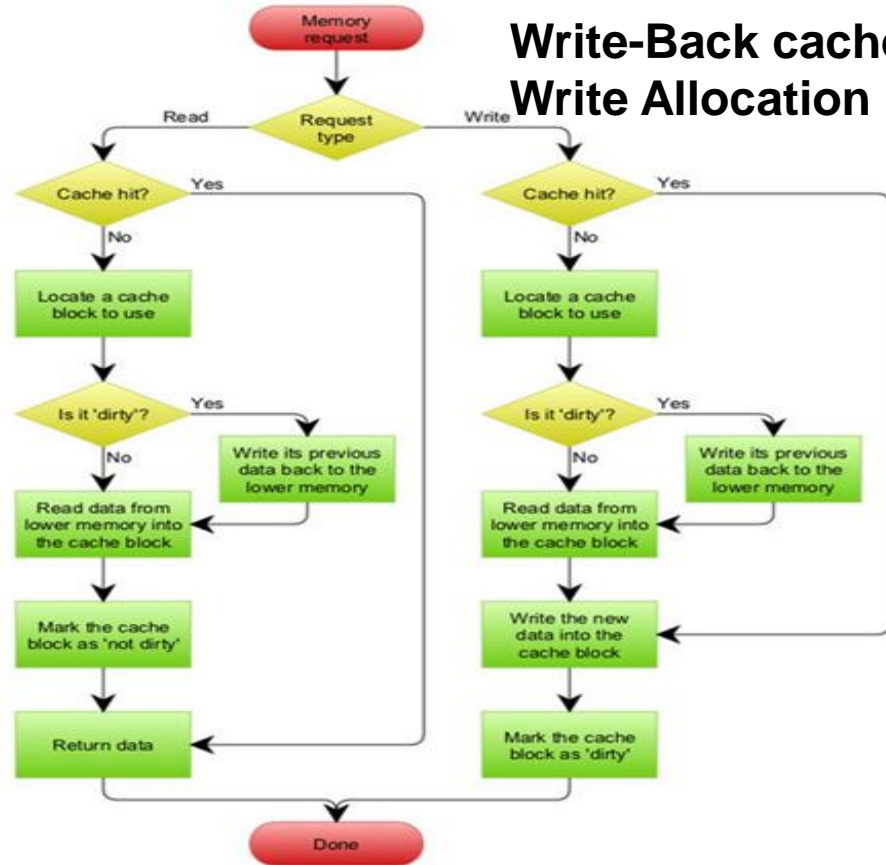


Write strategy

Write-Through cache with No-Write Allocation



Write-Back cache with Write Allocation



Accessing Cache Memory



$$\text{Average memory access time} = \text{Hit time} + (\text{Miss rate} \times \text{Miss penalty})$$

- ❖ **Hit Time:** Time to find the block in the cache and return it to processor [*indexing, tag comparison, transfer*].
- ❖ **Miss Rate:** Fraction of cache access that result in a miss.
- ❖ **Miss Penalty:** Number of additional cycles required upon encountering a miss to fetch a block from the next level of memory hierarchy.

How to optimize cache ?

- ❖ Reduce Average Memory Access Time
- ❖ $AMAT = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$
- ❖ Motives
 - ❖ Reducing the miss rate
 - ❖ Reducing the miss penalty
 - ❖ Reducing the hit time

Larger Block Size

- ❖ **Larger block size to reduce miss rate**


- ❖ **Advantages**

- ❖ Utilize spatial locality
- ❖ Reduces compulsory misses

- ❖ **Disadvantages**

- ❖ Increases miss penalty
- ❖ More time to fetch a block to the cache [bus width issue]
- ❖ Increases conflict misses
- ❖ More number of blocks will be mapped to the same location
- ❖ May bring useless data and evict useful data [pollution]

Larger Block Size



Larger Caches

- ❖ **Larger cache to reduce miss rate**

- ❖ **Advantages**

 - ❖ Reduces capacity misses

 - ❖ Can accommodate larger memory footprint

| Block size | Cache size | | | |
|------------|------------|-------|-------|-------|
| | 4K | 16K | 64K | 256K |
| 16 | 8.57% | 3.94% | 2.04% | 1.09% |
| 32 | 7.24% | 2.87% | 1.35% | 0.70% |
| 64 | 7.00% | 2.64% | 1.06% | 0.51% |
| 128 | 7.78% | 2.77% | 1.02% | 0.49% |
| 256 | 9.51% | 3.29% | 1.15% | 0.49% |

- ❖ **Drawbacks**

 - ❖ Longer hit time

 - ❖ Higher cost, area and power

Larger Caches



Higher Associativity

❖ Higher associativity to reduce miss rate

- ❖ Fully associative caches are the best, but high hit time.
- ❖ So increase the associativity to an optimal possible level

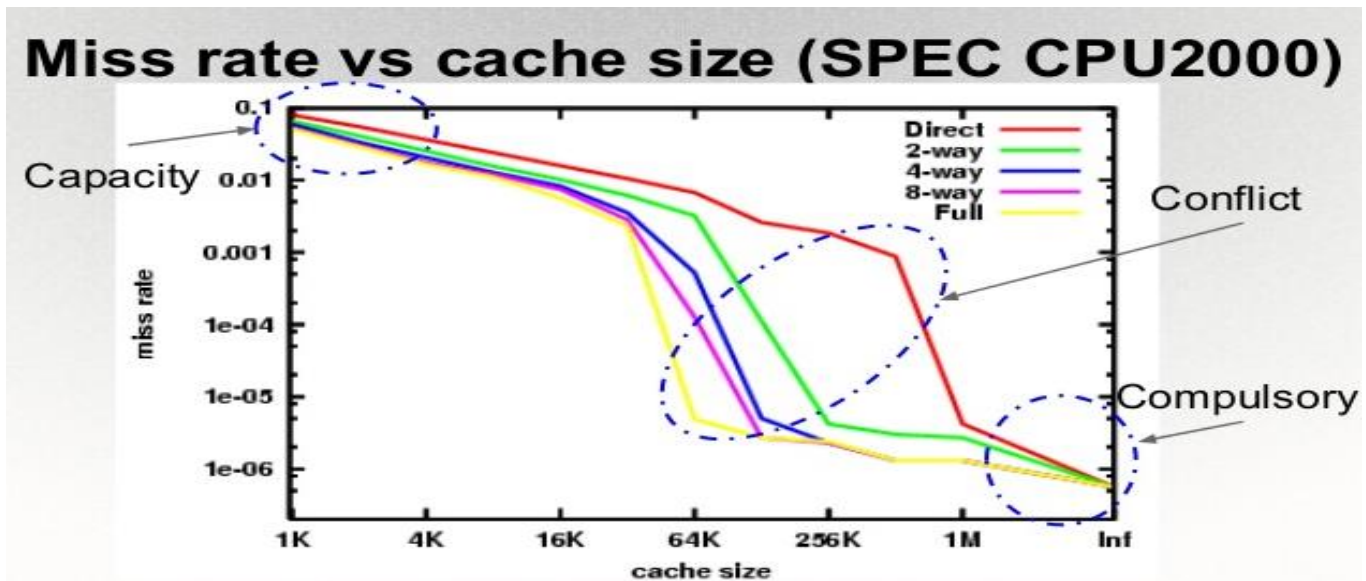
❖ Advantages

- ❖ Reduce conflict miss
- ❖ Reduce miss rate and eviction rate

❖ Drawbacks

- ❖ Increase in the hit time
- ❖ Complex design than direct mapped
- ❖ More time to search in the set (tag comparison time)

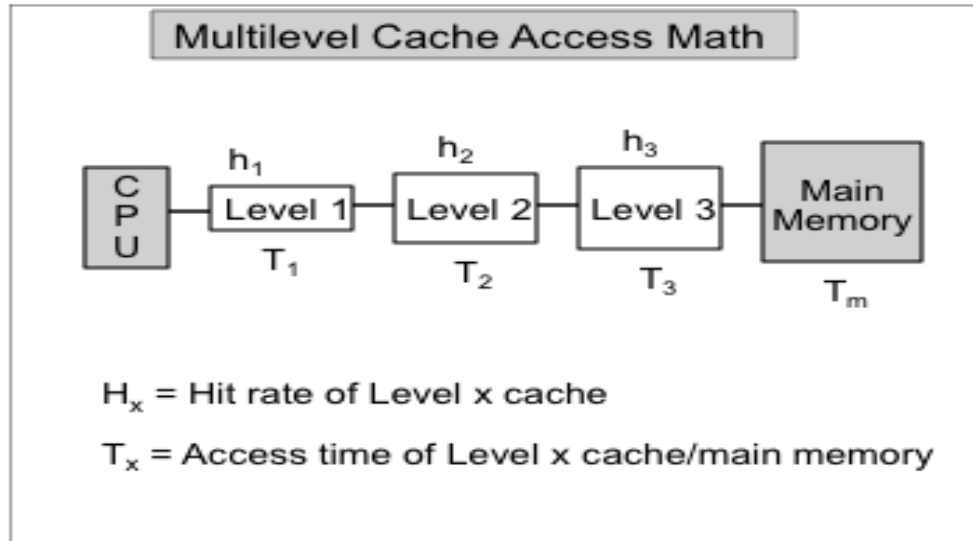
AMAT vs cache associativity



Multilevel caches

- ❖ **Multilevel caches to reduce miss penalty**
- ❖ **Caches should be faster** to keep pace with the speed of processors, **AND cache should be larger** to overcome the widening gap between the processor and main memory
- ❖ Add another level of cache between the cache and memory.
- ❖ The first-level cache (L1) can be small enough to match the clock cycle time of the fast processor. [Low hit time]
- ❖ The second-level cache (L2) can be large enough to capture many accesses that would go to main memory, thereby lessening the effective miss penalty. [Low miss rate]

Multilevel caches



Average memory access time = Hit time_{L1} + Miss rate_{L1} × Miss penalty_{L1}

Miss penalty_{L1} = Hit time_{L2} + Miss rate_{L2} × Miss penalty_{L2}

Average memory access time = Hit time_{L1} + Miss rate_{L1}
× (Hit time_{L2} + Miss rate_{L2} × Miss penalty_{L2})

Reference

- ❖ **Computer Architecture-A Quantitative Approach** (5th edition),
John L. Hennessy, David A. Patterson, Morgan Kaufman.
- ❖ Chapter 2: Memory Hierarchy Design
 - ❖ Section 2.1: Introduction
- ❖ Appendix B: Review of Memory Hierarchy
 - ❖ Section B.2: Cache Performance
 - ❖ Section B.3: Basic Cache Optimizations
- ❖ **NPTEL Video Link:** <https://tinyurl.com/yf94dnby>



johnjose@iitg.ac.in
<http://www.iitg.ac.in/johnjose/>