

CS223 : Computer Architecture & Organization

Lecture 22 [31.03.2022]

Advanced Cache Optimizations



Dr. John Jose

Associate Professor

**Department of Computer Science & Engineering
Indian Institute of Technology Guwahati, Assam.**

Accessing Cache Memory



$$\text{Average memory access time} = \text{Hit time} + (\text{Miss rate} \times \text{Miss penalty})$$

- ❖ **Hit Time:** Time to find the block in the cache and return it to processor [*indexing, tag comparison, transfer*].
- ❖ **Miss Rate:** Fraction of cache access that result in a miss.
- ❖ **Miss Penalty:** Number of additional cycles required upon encountering a miss to fetch a block from the next level of memory hierarchy.


How to optimize cache ?

- ❖ Reduce Average Memory Access Time
- ❖ $AMAT = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$
- ❖ Motives
 - ❖ Reducing the miss rate
 - ❖ Reducing the miss penalty
 - ❖ Reducing the hit time

Multilevel caches

- ❖ **Multilevel caches to reduce miss penalty**
- ❖ **Caches should be faster** to keep pace with the speed of processors, **AND cache should be larger** to overcome the widening gap between the processor and main memory
- ❖ Add another level of cache between the cache and memory.
- ❖ The first-level cache (L1) can be small enough to match the clock cycle time of the fast processor. [Low hit time]
- ❖ The second-level cache (L2) can be large enough to capture many accesses that would go to main memory, thereby lessening the effective miss penalty. [Low miss rate]

Multilevel caches



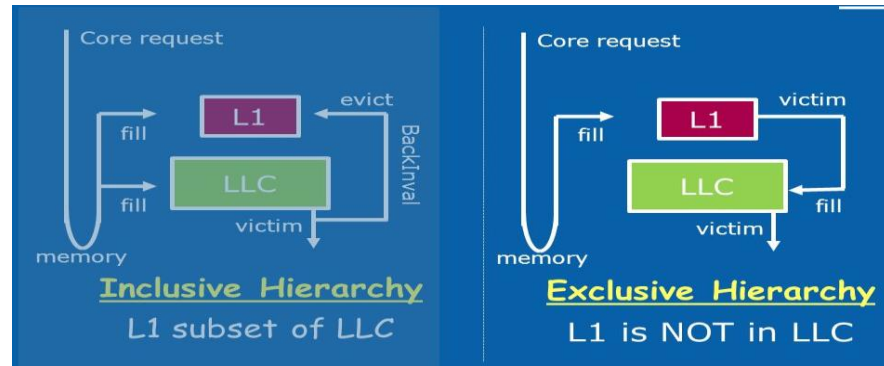
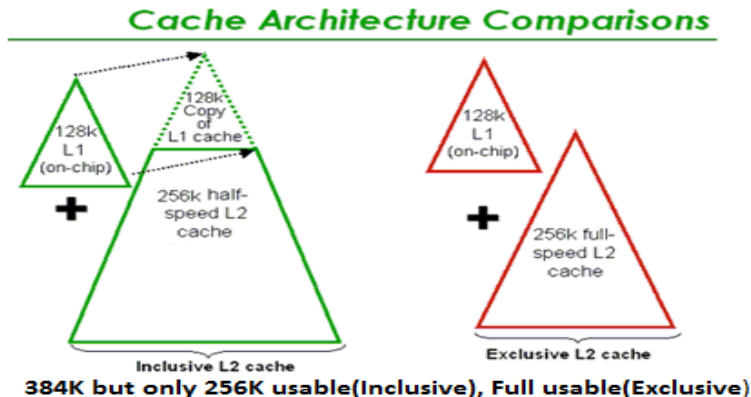
Average memory access time = Hit time_{L1} + Miss rate_{L1} × Miss penalty_{L1}

Miss penalty_{L1} = Hit time_{L2} + Miss rate_{L2} × Miss penalty_{L2}

Average memory access time = Hit time_{L1} + Miss rate_{L1}
× (Hit time_{L2} + Miss rate_{L2} × Miss penalty_{L2})

Multilevel caches

- ❖ **Multilevel caches to reduce miss penalty**
- ❖ **Local miss rate:** Number of misses in a cache level divided by number of memory access to this level.
- ❖ **Global miss rate:** Number of misses in a cache level divided by number of memory access generated by the CPU.
- ❖ **Inclusive and Exclusive caches**



Prioritize read miss over writes

- ❖ **Prioritize read misses to reduce miss penalty**
- ❖ If a read miss has to evict a dirty memory block, the normal sequence is write the dirty block to memory and read the missed block


```
SW R3, 512(R0)      ;M[512] ← R3      (cache index 0)
LW R1, 1024(R0)     ;R1 ← M[1024]    (cache index 0)
LW R2, 512(R0)      ;R2 ← M[512]    (cache index 0)
```

- ❖ Optimization: copy the dirty block to a buffer, read from memory and then write the block - reduces CPU's waiting time on read miss


Way Prediction

- ❖ **Predict the way in a set to reduce hit time**
- ❖ To improve hit time, predict the way to pre-set multiplexer
- ❖ Extra bits are set to predict the block with in the set
- ❖ Remember the MRU way of a given set
- ❖ Using the prediction bits, power gating can be done on unused ways for reducing power.

Way Prediction




Way Prediction - Correct



Way Prediction Miss

Victim Caches

- ❖ **Introduce victim caches to reduce mis- penalty**
- ❖ Additional storage near L1 for MR evicted blocks
- ❖ Efficient for thrashing problem in L1 cache
- ❖ Look up Victim Caches before L2
- ❖ L1 and Victim caches are exclusive




Pipelining Cache

- ❖ **Pipelined cache for faster clock cycle time.**
- ❖ Split cache memory access into several sub-stages
- ❖ Ex: Indexing, Tag Read, Hit/Miss Check, Data Transfer
- ❖ Pipeline cache access to improve bandwidth
 - ❖ Examples: Pentium 4 – Core i7: 4 cycles
- ❖ But slow in hits.
- ❖ Increases branch mis-prediction penalty
- ❖ Makes it easier to increase associativity

Multi-banked Caches


- ❖ **Multi-banked caches to increase cache bandwidth**
- ❖ Rather having a single monolithic block, divide cache into independent banks that can support simultaneous accesses.
 - ❖ ARM Cortex-A8 supports 1-4 banks for L2
 - ❖ Intel i7 supports 4 banks for L1 and 8 banks for L2
- ❖ Sequential Interleaving



Four-way interleaved cache banks using block addressing. Assuming 64 bytes per blocks, each of these addresses would be multiplied by 64 to get byte addressing.

Non-blocking Caches

- ❖ **Non blocking caches to increase cache bandwidth**
- ❖ Caches can serve hits under multiple cache miss in progress
 - ❖ *(a) Hit under miss (b) Hit under multiple miss*



- ❖ Must needed for OOO superscalar processor for IPC increase
- ❖ L2 must support it with L1-MSHR (Miss Status Holding Reg.)
- ❖ On an L1 miss allocate MSHR entry, clear upon L2 respond with cache block reply.

Non-blocking Caches

- ❖ **Non blocking caches to increase cache bandwidth**
- ❖ Processors can hide L1 miss penalty but not L2 miss penalty
- ❖ Reduces the effective miss penalty by overlapping miss latencies
- ❖ Significantly increases the complexity of the cache controller as there can be multiple outstanding memory accesses
- ❖ Requires pipelined or banked memory system

Early Restart

- ❖ **Early restart to reduce miss penalty**
- ❖ Do not wait for entire block to be loaded for restarting CPU
- ❖ **Early restart**
 - ❖ Request words in normal order
 - ❖ Send missed work to the processor as soon as it arrives
 - ❖ L2 controller is not involved in this technique

- **Early restart**



1 -> 2 -> 3 -> 4

Requested word: word 3



Processing performed
background

Critical Word First

- ❖ **Critical word first to reduce miss penalty**

- ❖ **Critical word first**

- ❖ Request missed word from memory first
- ❖ Send it to the processor as soon as it arrives
- ❖ Processor resume while rest of the block is filled in cache
- ❖ L2 cache controller forward words of a block out of order.
- ❖ L1 cache controller should re-arrange words in block
 - **Critical word first**



3 -> 4 -> 1 -> 2

As soon as 3 is brought it is forwarded to CPU
In background 4, 1 and 2 is brought

Merging Write Buffer

- ❖ **Write buffer merging to reduce miss penalty**
- ❖ Write buffer allows the processor to continue without waiting for writes to get over.
- ❖ When performing a store/write on a block that is already pending in the write buffer, update write buffer
- ❖ Reduces stalls due to full write buffer, improve buffer efficiency
- ❖ If buffer is full writes incur processor stall.

Write address	V		V		V		V	
100	1	Mem[100]	0		0		0	
108	1	Mem[108]	0		0		0	
116	1	Mem[116]	0		0		0	
124	1	Mem[124]	0		0		0	

•Each 64-bit word takes up one entry

No write buffering

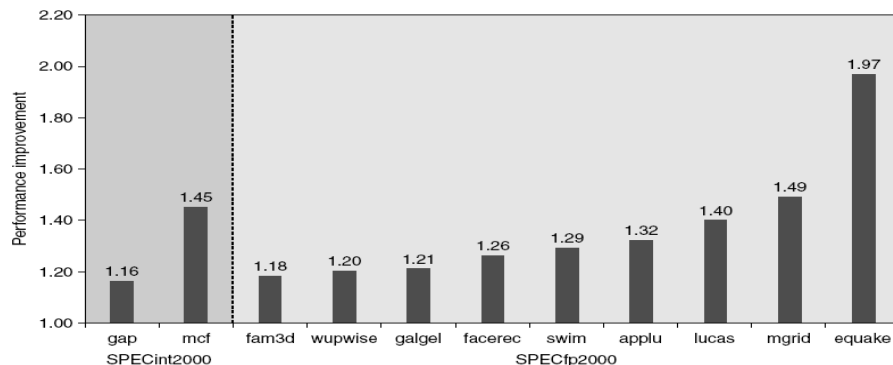
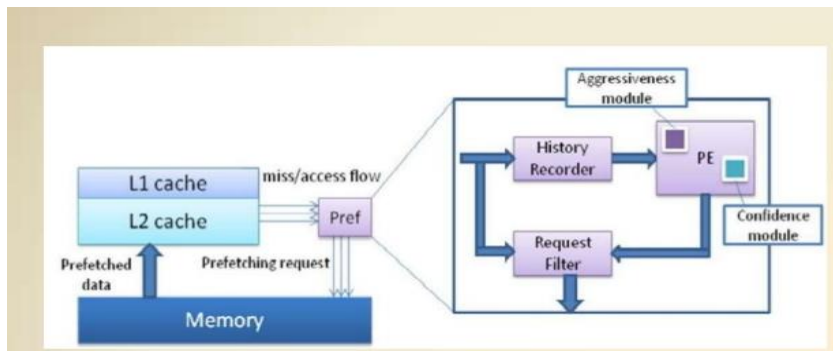
Write address	V		V		V		V	
100	1	Mem[100]	1	Mem[108]	1	Mem[116]	1	Mem[124]
	0		0		0		0	
	0		0		0		0	
	0		0		0		0	

•Sequential addresses can be placed in one entry

Write buffering

Hardware Prefetching

- ❖ **Pre-fetching to reduce miss rate and miss penalty.**
- ❖ Pre-fetch items before processor request them.
- ❖ Fetch more blocks on miss -include next sequential block
- ❖ Requested block is kept in I-cache and next in stream buffer.
- ❖ If a missed block is in stream buffer, cache miss is cancelled



Compiler Controlled Pre-fetching

- ❖ **Pre-fetching to reduce miss rate and miss penalty.**
- ❖ Insert pre-fetch instructions before data is needed
- ❖ Pre-fetching will give performance only if processor reads from caches and executes while pre-fetching is in progress.
- ❖ Register pre-fetch
 - ❖ Loads data into register
- ❖ Cache pre-fetch
 - ❖ Loads data into cache
- ❖ Use loop unrolling and scheduling for pre-fetching data of adjacent iterations

Reference

- ❖ **Computer Architecture-A Quantitative Approach** (5th edition),
John L. Hennessy, David A. Patterson, Morgan Kaufman.
- ❖ Chapter 2: Memory Hierarchy Design
 - ❖ Section 2.1: Introduction
 - ❖ Section 2.2: Ten Advanced Cache Optimizations of Cache Performance
- ❖ **NPTEL Video Links:**
 - ❖ <https://tinyurl.com/yf94dnby>
 - ❖ <https://tinyurl.com/yfgmfsmm>



johnjose@iitg.ac.in
<http://www.iitg.ac.in/johnjose/>